ATHARVA GANGRADE

Charlotte, NC 28262 | (704) 369-6432 | atharvagangrade8@gmail.com | linkedin.com/in/atharva-gangrade | https://agangrad2.wixsite.com/atharva-gangrade

PROFILE SUMMARY

Computer Science student building production-grade ML systems and distributed data pipelines processing large-scale datasets. Improved model accuracy (R^2 up to 0.89) and reduced system runtime by up to 28% through optimization and scalable backend design. Strong in Python, Java, distributed systems, and performance-driven engineering.

EDUCATION

University of North Carolina at Charlotte — **Charlotte, NC**
B.S. Computer Science — *August 2023 – May 2027(Expected)*

- **GPA:** 3.66
- **Dean's List :** Fall 2023, Spring 2024
- **Chancellor's List :** Fall 2024, Spring 2025, Fall 2025
- **Relevant Courses:** Introduction to Java, Introduction to Python, Introduction to Software Engineering, Data Structures and Algorithms, Database Design and Implementation, Introduction to Operating Systems and Networking, Data Mining, Cloud Computing and Data Analysis, Introduction to Machine Learning, Visual Analytics

TECHNICAL SKILLS

Programming: Python, Java, C#, SQL

ML & Data: Pandas, NumPy, Scikit-Learn, Random Forest, Regression, Feature Engineering

Distributed Systems: Hadoop MapReduce, Spark

Backend & APIs: Django, React, REST APIs

Cloud & DevOps: AWS, Azure, Docker

Data Tools: Power BI, Git

PROJECTS

Global Food Wastage Predictive Analytics | End-to-End ML Pipeline — *Oct 2025 – Dec 2025*
GitHub link : https://github.com/atharva-0085/Global-Food-Wastage-Predictive-Analytics

- Built end-to-end ML pipeline on multi-country, multi-year economic dataset (feature engineering, model tuning, validation).
- Optimized Random Forest (400 estimators), achieving $R^2 = 0.87$ and reducing RMSE by 18% vs baseline on 20% holdout.
- Identified top economic loss drivers via feature importance and correlation modeling to guide data-backed risk analysis.
- Designed reproducible training workflow with structured experimentation and evaluation tracking.

Used Car Price Prediction | Advanced Regression Modeling & Feature Intelligence — *June 2025 – Aug 2025*
GitHub link : https://github.com/atharva-0085/Used-Car-Price-Prediction

- Engineered structured dataset (encoding, scaling, outlier handling), improving 5-fold CV performance by 13%.
- Benchmarked regression models, selecting optimal model achieving $R^2 = 0.89$ with minimized RMSE/MAE.
- Extracted key pricing drivers to support data-driven valuation strategy.
- Applied systematic hyperparameter tuning for performance optimization.

Chicago Food Inspections Big Data Pipeline | Hadoop MapReduce at Scale — *Aug 2025 – Dec 2025*
GitHub link : https://github.com/atharva-0085/Chicago-Food-Inspections-

- Engineered distributed Hadoop MapReduce jobs in Java to process large public datasets for compliance and risk aggregation.
- Reduced end-to-end runtime by 28% through partitioning strategy, combiner optimization, and reducer parallelism.
- Implemented fault-tolerant batch workflows for reliable large-scale data processing.
- Optimized key distribution to minimize shuffle bottlenecks.

UniSphere | Full-Stack Platform for International Students — *Feb 2025 – May 2025*
GitHub link : https://github.com/atharva-0085/UniSphere

- Architected scalable Django backend with normalized relational schema and REST APIs supporting housing & roommate workflows.
- Reduced query latency by 15% via indexing and execution plan optimization under concurrent usage.
- Implemented modular backend architecture with secure authentication and role-based access control.
- Designed database schema to minimize redundancy and improve join efficiency.